

Contact Information	739 E El Camino Real Apt 112 Sunnyvale, CA-94087	Email:cvkkumar@cs.stanford.edu Ph: 650-384-9476
Interests	Large Scale Data Mining and Analysis, Machine Learning, Information Retrieval and Software Engineering	
Education	Stanford University, M.S. Computer Science	Sep '08 - Apr '10 GPA : 3.86 / 4.0
	<ul style="list-style-type: none">• Recipient of the prestigious <i>BPCL Scholarship</i> awarded by Bharat Petroleum Corporation Ltd. awarded to the elite few for graduate study in the US.• <i>Relevant Coursework:</i> Machine Learning, Data Mining & E-Business, Modern Applied Statistical Learning, Computational Advertising, Information Retrieval & Web Search, Probabilistic Graphical Models(A), Data Mining, Transaction Processing and Distributed Databases, Network Analysis.	
	Birla Institute of Technology and Science, Pilani - Goa Campus M Sc.(Tech.) Information Systems	Aug 04 - Jun 08 CGPA : 9.95/10.0
Experience	Apple Inc.	
	Software Engineer - Maps Data Insights	Apr '10 to present May '13 - present
	<ul style="list-style-type: none">• TODO	
	Software Engineer - Internet Services	Apr '10 - May '13
	Reputation and Indirect Fraud Detection on iTunes AppStore	
	<ul style="list-style-type: none">• Responsible for the end-to-end development of the fraud solution to reputation fraud on iTunes AppStore as a part of a very small team - right from requirements gathering, conception, design, implementation and evaluation.• Experienced with data mining, network analysis and machine learning on massive scales.• Implemented a scalable version of the <i>Belief propagation algorithm</i> that works on graphs with billion nodes and tens of billions of edges.• Also responsible for communicating the working of system to different stakeholders.• System currently in production for more than a year and has a very high hit-rate.• Used Pig, Java MapReduce, Mahout and HBase on the Hadoop ecosystem.	
	Account Creation Fraud Detection on iTunes	
	<ul style="list-style-type: none">• Designed and implemented a high performance and scalable subsystem of a system that risks account creations that works in the path of the transaction.• Implemented a scalable stream mining algorithm that works on continuous data streams and still handling hundreds of transactions per second.• Implementation was carried out in Java.	
	Spam Detection on iTunes Ping comments	
	<ul style="list-style-type: none">• Worked on machine learning techniques for spam detection on Ping.	
Summer Intern		June '09 to Mar '10
	<ul style="list-style-type: none">• Worked on automated sentiment analysis and opinion mining from Micro-blogs using a combination of sophisticated machine learning and data mining techniques.• Part of a two-member team that was responsible for the entire project from its conception and design to implementation and production deployment.• Project involved use and analysis of several state-of-the-art machine learning techniques for feature engineering, feature selection, skew handling in datasets, model comparisons for supervised learning, ensemble techniques for classifiers and evaluation of results.• Used Java, libSVM, Weka.	
Hewlett Packard Labs-India <i>STAIR : System for Topical and Aggregated Information Retrieval.</i>		Research Intern Jan 08 to June 08

- Developed the architecture and the prototype of STAIR - an IR system that applied of a combination of Collaborative analysis and Focused Crawling techniques on the web documents to provide personalized, consolidated information relevant to the used as an aggregated PDF document.
- Implemented using Java, on top of Lucene. The semantics information were obtained from WordNet.
- Published as a HP Labs Technical Report in 2009.

Center for Study of Language and Information, Stanford University

Graduate Research Assistant

Aug 08 to June 09

Cognitive Assistant that Learns and Organizes (CALO)

- Member of a team working on CALO, a system that extracts decisions from multi-party meetings to enable the effective handling of feedbacks.
- Also involved in the evaluation of new features to decision extraction process.
- Implementation used Java, Swing and libSVM.

Bhabha Atomic Research Center, Trombay Project Intern at DRHR

May 06 to July 06

Image Processing and Software Development for Simplifying Robot Trajectory Generation

- Built a system that extracts information of a continuous path from any arbitrary raw image using graph-theoretic methods and provides input to the indigenous *Sensor-cum-Manipulator*, a Parallel planar kinematic robot.
- Implemented using C, Matlab (for image processing) and VB (as a wrapper GUI).

Selected Projects

Role Discovery in Social Networks using Dirichlet Multinomial Regression Based Topic Modeling

- Implemented an unsupervised algorithm for identification of hierarchical roles in an organization using a combination of Social Network Analysis, Spectral Clustering and a variant of LDA (Latent Dirichlet Allocation).
- *Instructor: Prof. Daphne Koller, Stanford University*

Recommendation Systems based on Delicious and Twitter

- Implemented a People Recommendation System on Twitter (in Python) by a combination of several algorithms, that included collaborative filtering and network analysis.
- Designed and implemented an URL Recommendation system by analysis of tags from *Delicious*.

URL Recommendation Based on Asymmetric Tag Similarity and Diffusion-Based Grouping

- Implemented an URL Recommendation system by analysis of tag similarity using data from ShareThis and Delicious using MapReduce and Partition based Joins on top of the Aster Cluster.

Finding Answerers on Yahoo! Answers

- Designed and implemented a system for selection of most appropriate answerers in *Yahoo! Answers* based on textual, structural and other auxiliary information. The result could be used to determine the routing for new questions.

Analysis of Text Based Classifiers

- Implemented and analyzed the performance of different Naïve Bayes classifiers on the 20-Newsgroups dataset, using Java and Lucene.

Comparison of Similarity Search Algorithms over Inverted Indexes

- Implemented and analyzed the performance of commonly used indexing similarity search algorithms - Term-at-a-Time and Document-at-a-Time. Also optimized the algorithms with efficient index compression.

RefMed - A Physician Referral and Review Service

- Designed and developed a physician referral and review service that enables patients to review and rate the physicians and facilitates the physicians to recommend other doctors to their patients.

An Ontology-based Automatic Staging system for Cancer

- Developed a Automatic Staging System for Breast Cancer over the NCI Thesaurus using SWRL, OWL and Protégé.

Implementation Of a Search Engine for Personalized Information Retrieval by profiling of user data

- Designed and implemented a prototype search engine that incorporates an additional dimension of Personalization through User Profiling for Enhanced relevancy.

Time Table Generator

- Implemented a system for automating the process of timetable generation for a University with its constraints, using a sub-optimal graph coloring approach for Constraint Satisfaction.

An Expert System for selection of Polymer Composite Systems

- Implemented a novel method for evaluation and ranking of constituent materials for composite products using *TOPSIS* , an MADM (Multiple Attribute Decision Making) approach that ensures an optimum solution for the characteristics desired.

Publications

C.V.Krishnakumar and Dr.Krishnan Ramanathan, *STAIR : A System for Topical and Aggregated Information Retrieval*, Proceedings of the International Conference on Intelligent Human Computer Interaction (IHCI) 2009.

Skills

Languages: Java, Pig, Python, C, SQL and basics of Hive.

Frameworks: Hadoop (Cloudera Certified Hadoop Developer), Spring, HBase, Mahout

Tools and Platforms: Weka, L^AT_EX, Basics of R and Teradata

Operating Systems: Unix based Systems, Microsoft Windows.

Honors and Achievements

Consistent recipient of the *Merit Scholarship* awarded by BITS Pilani to the top 10 students across the batch.

Secured the First Prize at *OpenSoft* - the software construction contest conducted as a part of QUARK-07, the national level technological fest at BITS-Pilani, Goa Campus.

Recipient of the Merit Certificate, awarded to the top 0.1% of students, for proficiency in English in AISSCE from the CBSE, 2004.